

Technical Computing for the Earth Sciences

EARS 80.03

Running Julia:

Locally, on your own computer: (recommended)

There are a few steps. First you need the language binary, and then you need an editor.

(1) Download a Julia binary for your OS from <https://julialang.org/downloads/>

(2) Pick an editor. Options include:

- a)** [VSCode, via Julia-vscode](#). Probably the default if you want a IDE-like experience, now that GitHub got bought by MS and stopped actively developing Atom. Follow instructions at: <https://github.com/julia-vscode/julia-vscode#installing-juliavs-codevs-code-julia-extension>
- b)** [Atom](#), via “Juno”. Follow directions at: <http://docs.junolab.org/latest/man/installation/>
- c)** A command-line editor, like Emacs or Vim. Potentially very efficient, but with steep learning curves. Directions at: <https://github.com/JuliaEditorSupport/julia-emacs> or <https://github.com/JuliaEditorSupport/julia-vim>
- d)** Editors that cost money, like Sublime Text or IntelliJ. Get plugins from <https://github.com/JuliaEditorSupport/Julia-sublime> or <https://github.com/JuliaEditorSupport/julia-intellij>

Remotely, on other people’s computers:

(i.e., “the cloud” — don’t have to install anything, but can be harder to get files in and out)

- a)** [repl.it](#) — Try starting at <https://julialang.org/learning/tryjulia/> or <https://repl.it/@logankilpatrick/TryJulia#main.jl>. Will need to make an account if you want to save changes, etc.
- b)** Binder notebook: e.g. <https://mybinder.org/v2/gh/binder-examples/demo-julia/master?filepath=demo.ipynb> (from <https://github.com/binder-examples/demo-julia>). Works pretty well if you like Jupyter notebooks, but times out if left inactive for more than 15 or 20 minutes. Great if you want to demonstrate your Julia code to other people.
- c)** Google Colab notebook: e.g. https://colab.research.google.com/github/ageron/julia_notebooks/blob/master/Julia_Colab_Notebook_Template.ipynb. Slightly inelegant (have to install Julia each time), but gives you access to free GPUs and TPUs (for now).

Julia basics

REPL MODES:

Normal: julia> type ? → help mode] → pkg mode ; → shell mode (e.g. BASH, etc.)

COMMON OPERATORS

$+$, $-$, $*$, $/$, \div , \wedge , $\%$, $\&$, $|$, $>$, $<$, \geq , \leq , \equiv , $|>$

exponent modulus OR pipe
 Integer division AND isequal
 $x|>\sin|>\cos$ is same as $\cos(\sin(x))$

SOME IMPORTANT TYPES

Numeric Types
 Float64 , Int64 , UInt64
 Int32 , 16 , (8) etc.*

*These are "primitive types".
 Can make more, but probably shouldn't

Arrays
 $\text{foo} = [1, 2, 3, 4, 5]$
 $\text{Array}\{\text{Float64}\}(\text{undef}, 100, 10)$

single quotes
 'a', 'b' - Characters
 "hello!" - String
 (like an array of characters)
 Double quotes

Tuples
 $a = (1, 2, 3, 4, 5)$
 Like arrays, but immutable

Composite Types
 struct Foo
 $a :: \text{Float64}$
 $b :: \text{Int32}$
 $c :: \text{Array}\{\text{Float64}, 1\}$
 end

Step ranges
 $a = 1:5$
 $b = 1:0.01:10$
 $c = 5:-1:1$
 Turn into an array with $\text{collect}(c)$
 Actually just composite types w/ methods for iteration

FUNCTIONS

```
function foo(a,b)
    c = 3*a + 7*b + 9
    return sin(c)
end
```

- Julia functions are built around multiple dispatch

- Dispatch favors function w/ most specific type signature

- All the operators above are just functions that can be extended like any other

- Can also write one-liners
 $\text{foo}(x) = \sin(x)$
 and anonymous functions
 $x \rightarrow 2x$

OTHER SYNTAX

Type annotations

$a :: \text{Float64}$

use in function signatures (for dispatch)
 & occasionally elsewhere (to avoid type instability)

"Dot-broadcasting"

$A .+ 5 .* \sin.(C)$

Applies the function/operator element-wise
 Can also write as
 $@. A + 5 * \sin(C)$

Julia basics

SOME USEFUL FUNCTIONS & MACROS*

- `typeof()`, `sizeof()`, `size()`, `bitstring()` - May help figure out what something is
- `@time`, `@allocated`, `@code_native`, `@code_warntype`, `@btine` - For measuring performance
- `print()`, `println()`, `Display()` - Show output in REPL

CONTROL FLOW

```
for i=1:100      i=0          if i<0  
    println(i)    while i<100    b=sqrt(-i)  
    end           i=i+1        else  
    end           println(i)   b=sqrt(i)  
    end           end
```

• Watch out for
global scope

GENERAL LANGUAGE STUFF

• Julia is "Just-in-time" compiled. This puts it somewhere in between
Compiled languages like Fortran, C, C++, Rust and Interpreted languages like Matlab, Python, Perl

It's not the only JIT-compiled language*, but is the only one where JIT compilation is used to allow efficient multiple dispatch

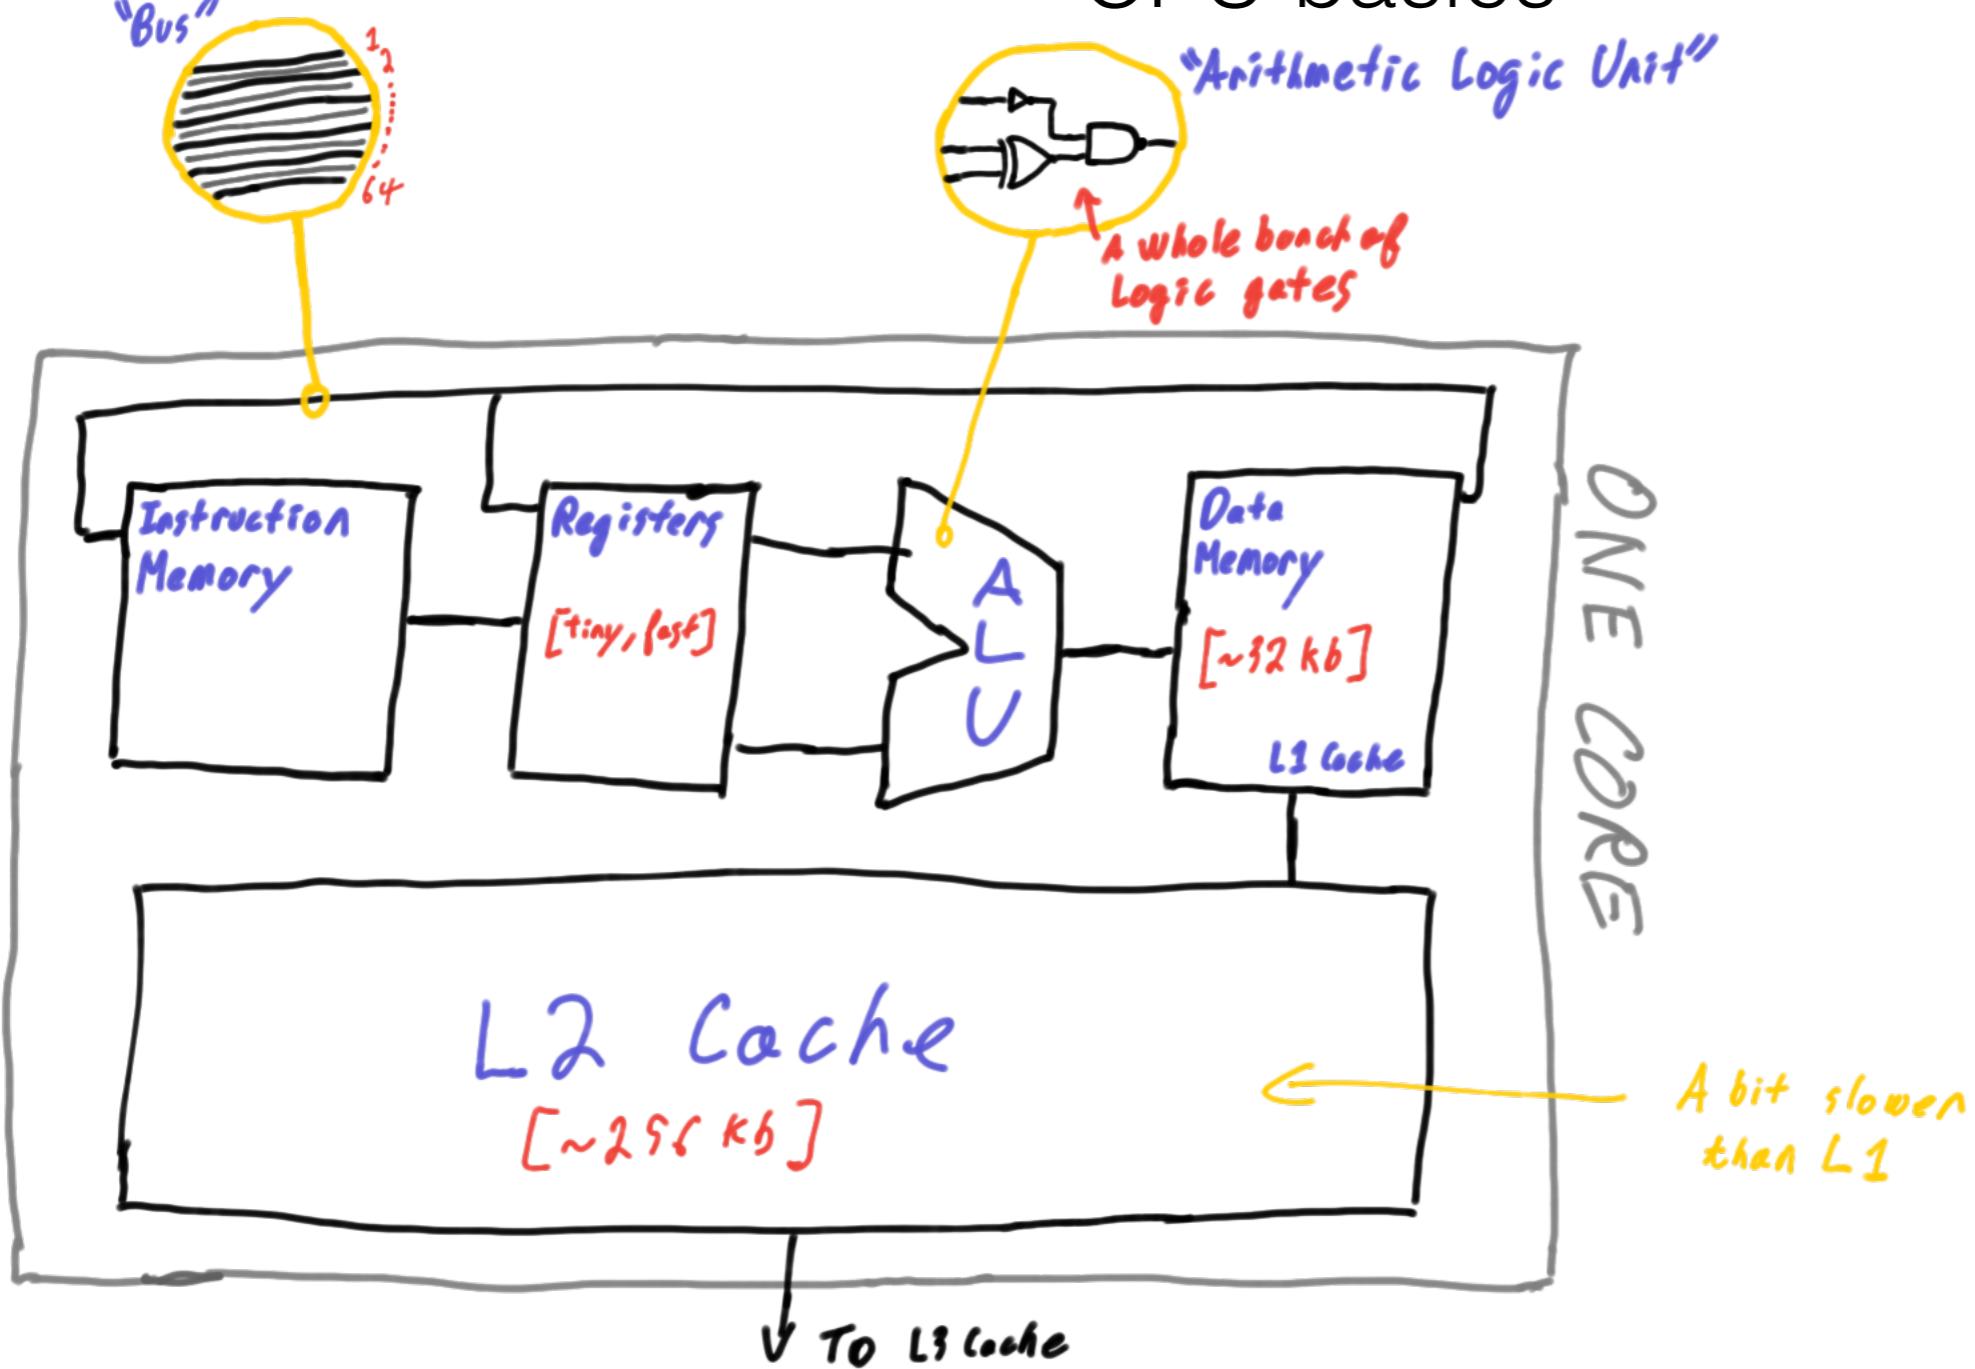


} Probably actually missing a bunch of steps
Only need to know if you want to do something really unconventional

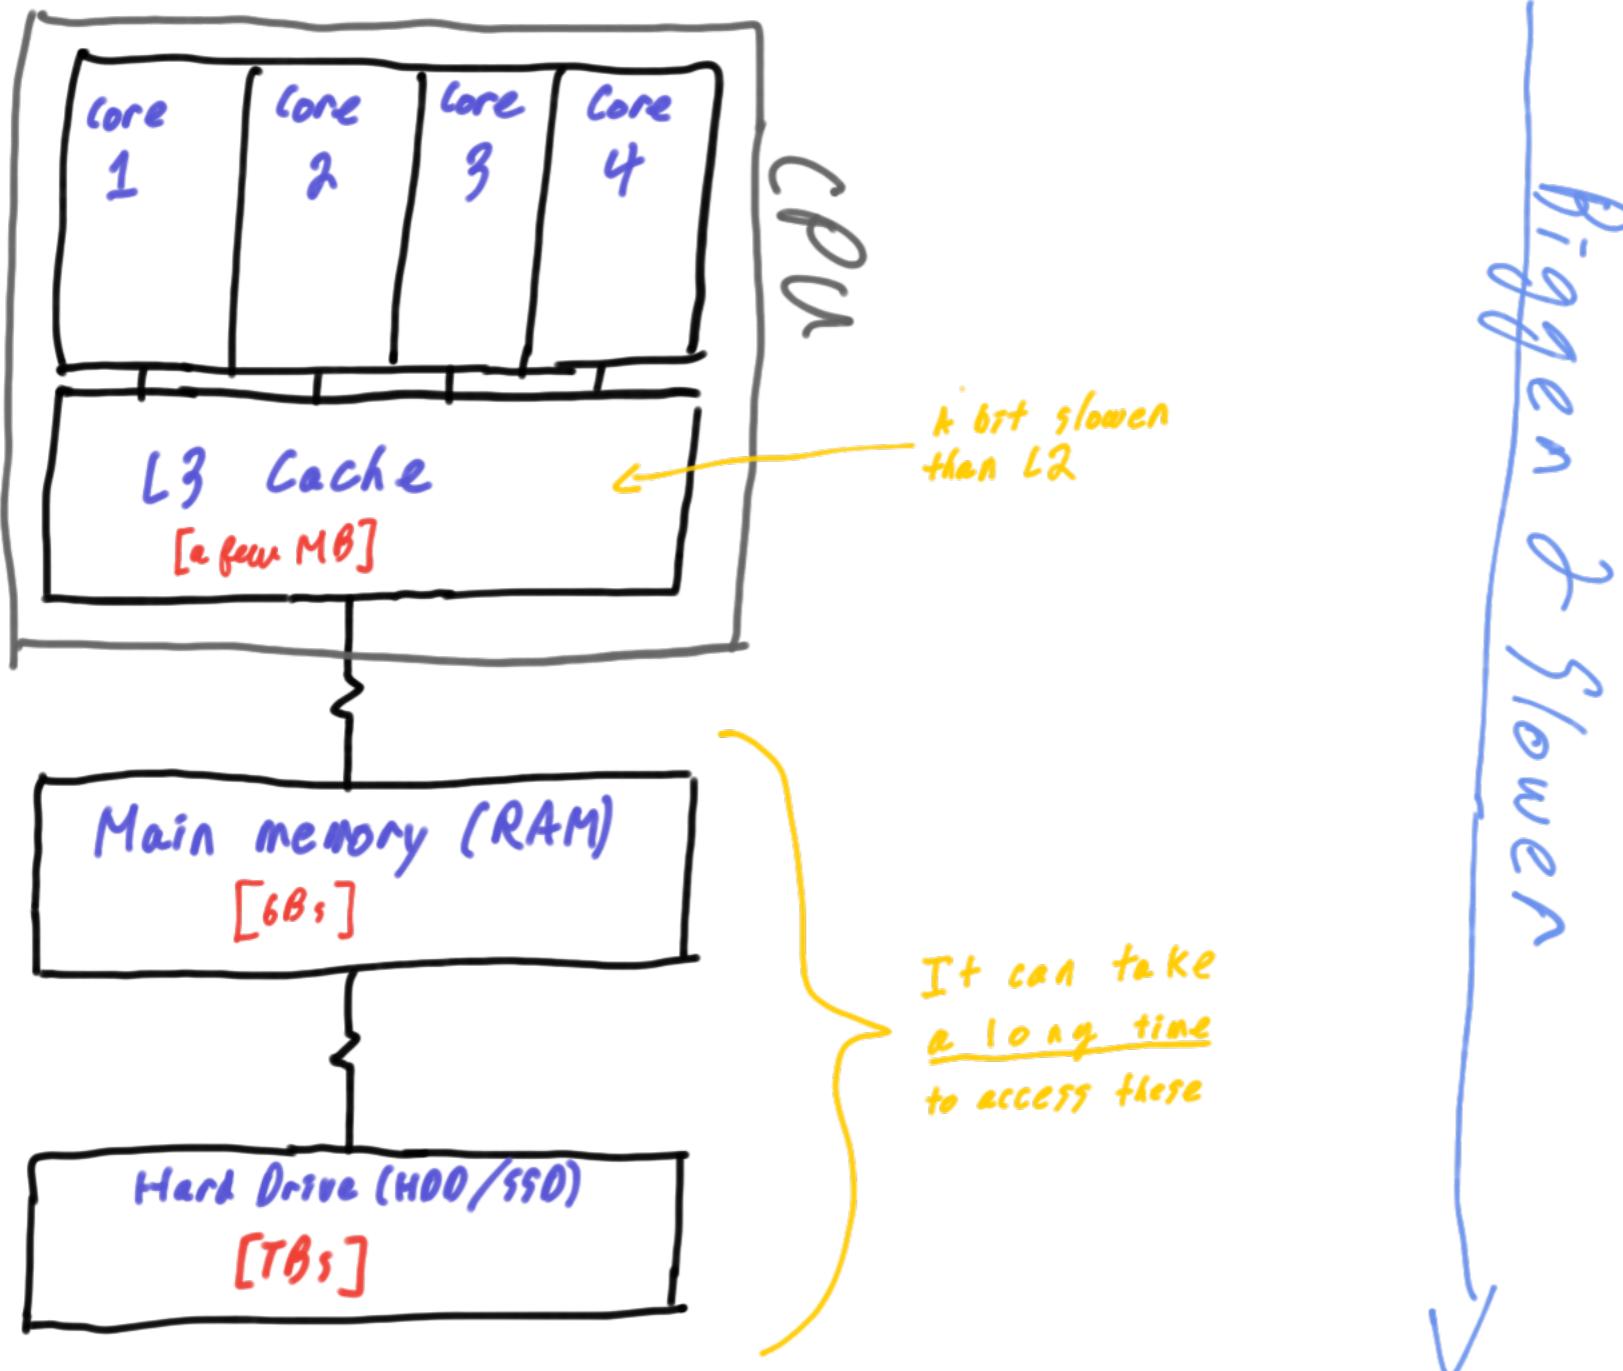
* JavaScript is often JIT compiled, because so much of the web depends on it. People have tried to JIT python, but the design of the language makes it hard.
Other languages are JITted to "bytecodes" that run on a VM (ex. Java) but not properly compiled to machine (assembly) code.

*Anything beginning with '@' is a macro, & that's all we're going to say about that for now

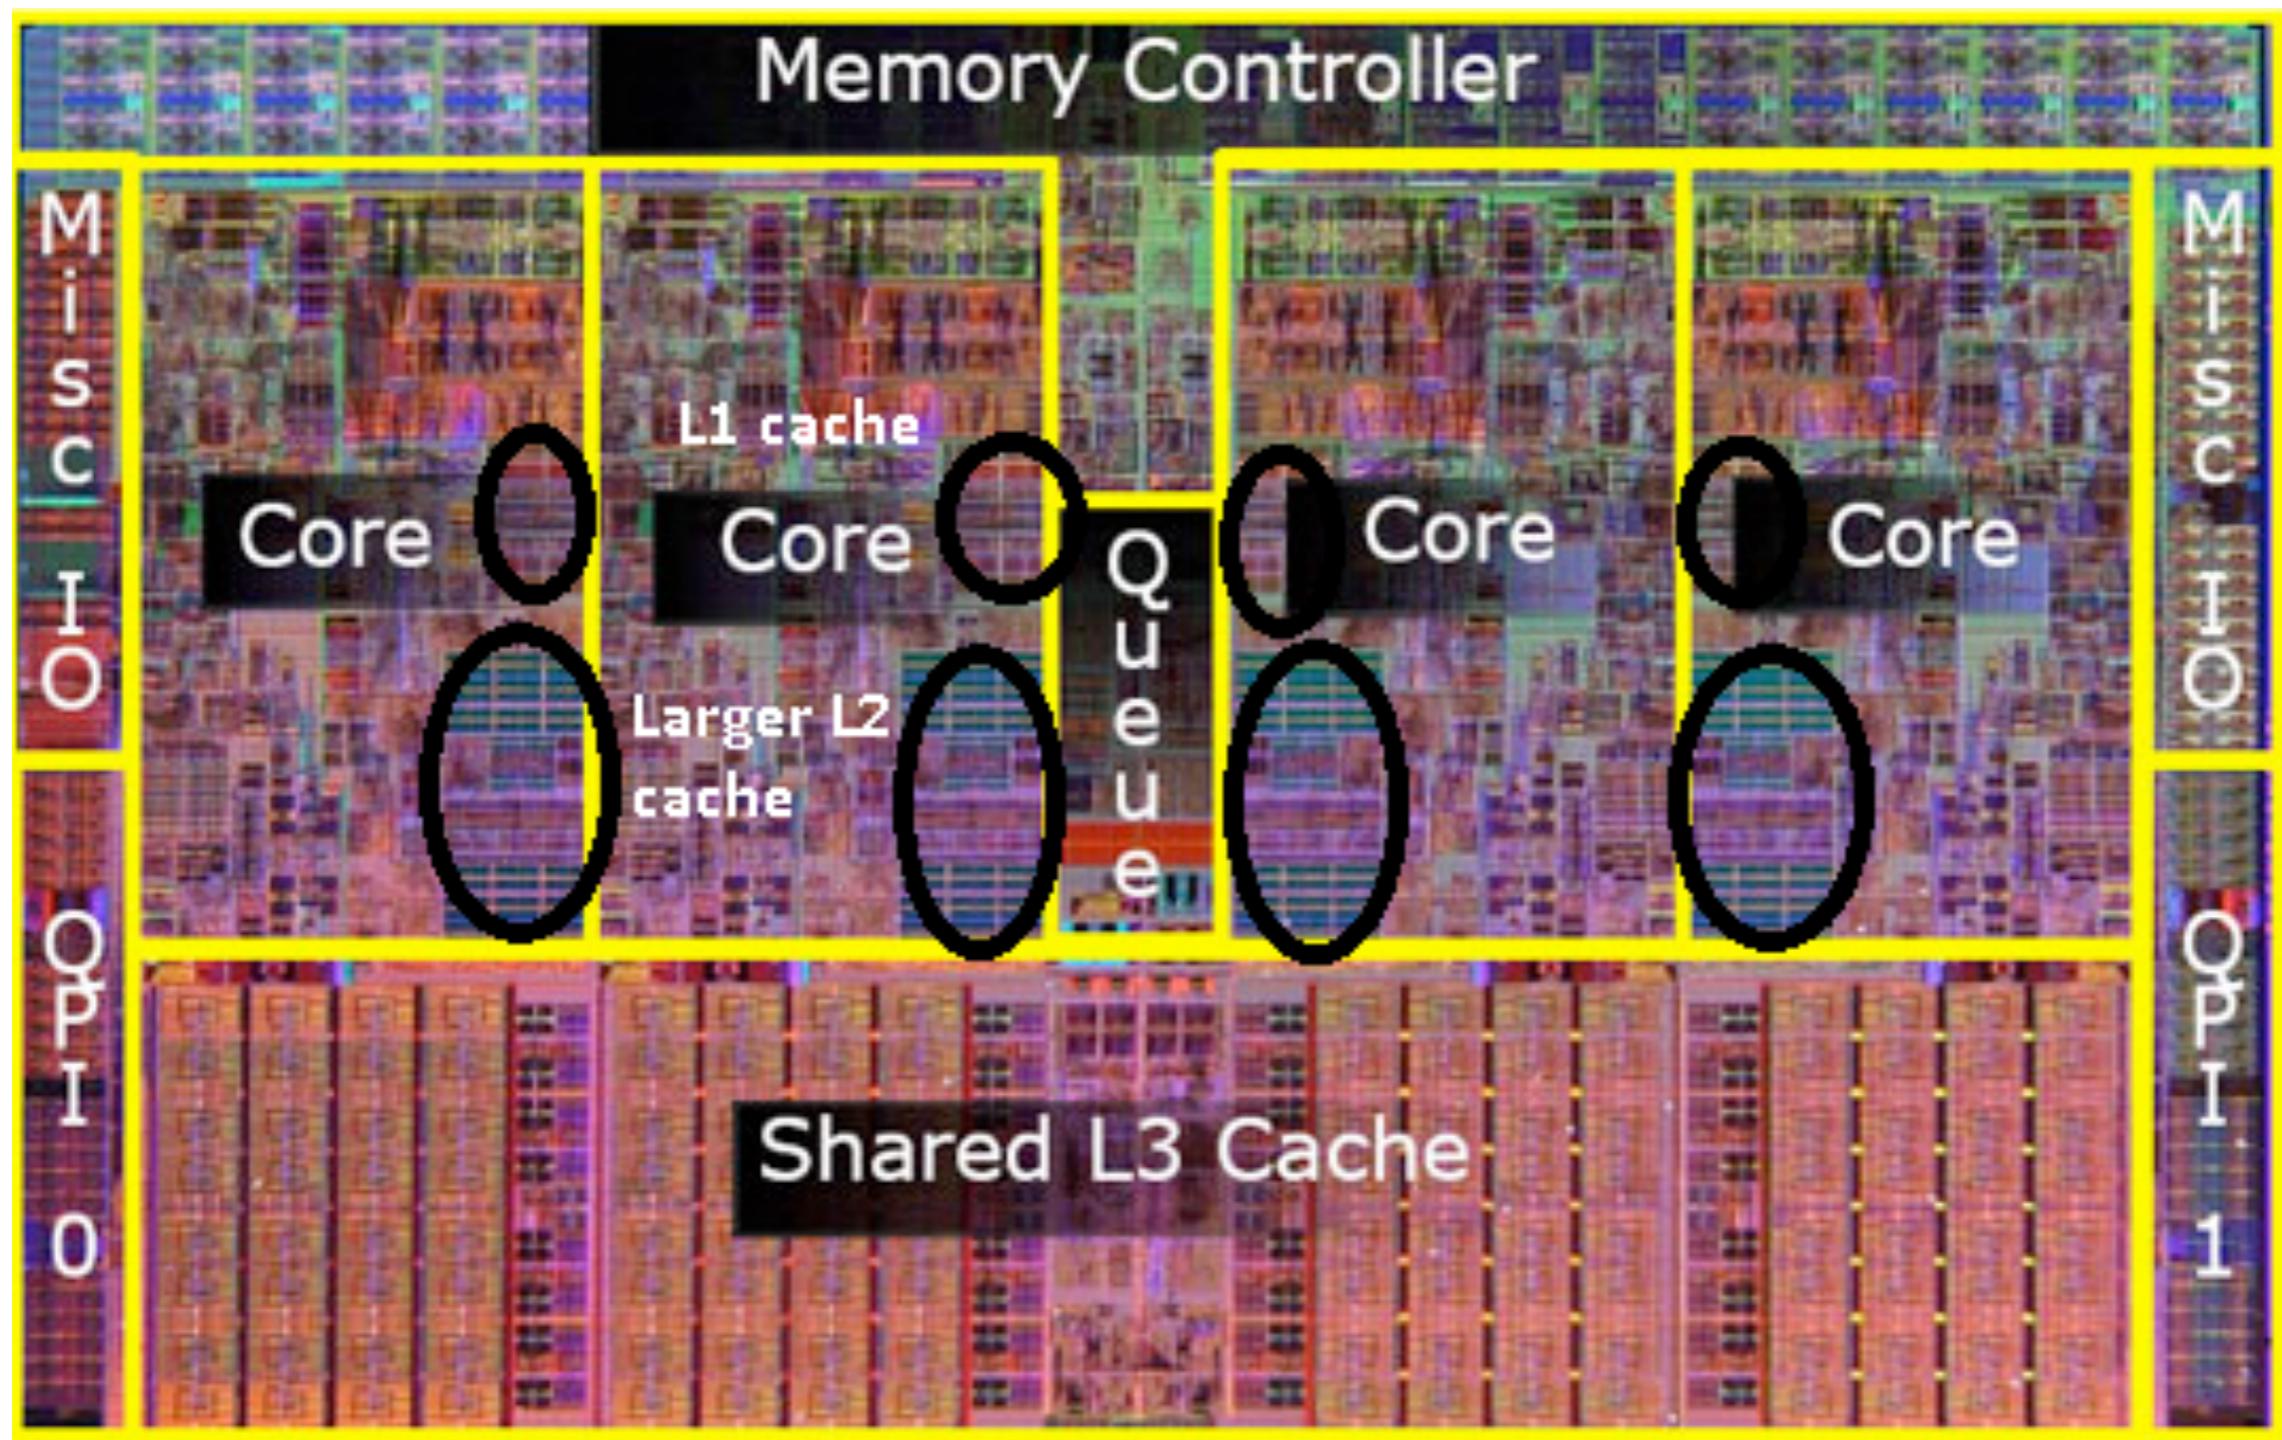
CPU basics



CPU basics



CPU basics



CPU basics

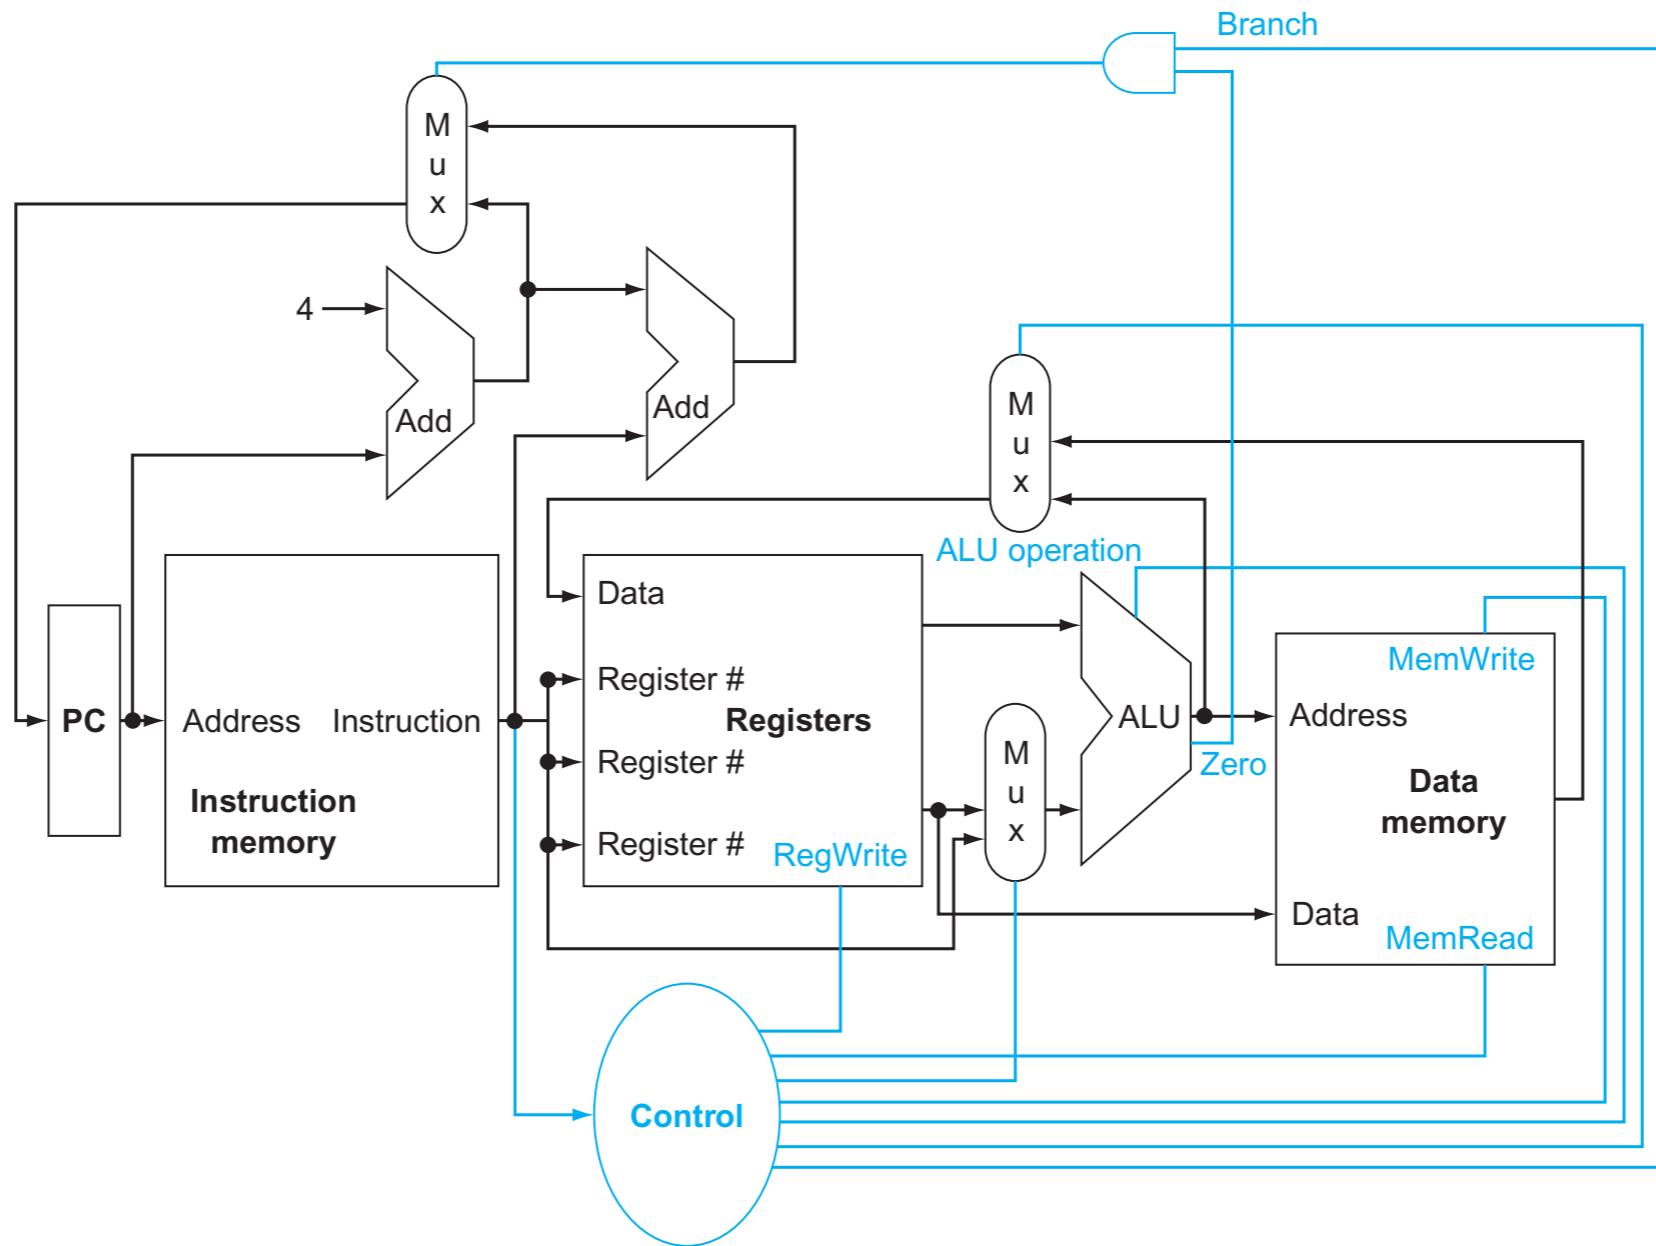
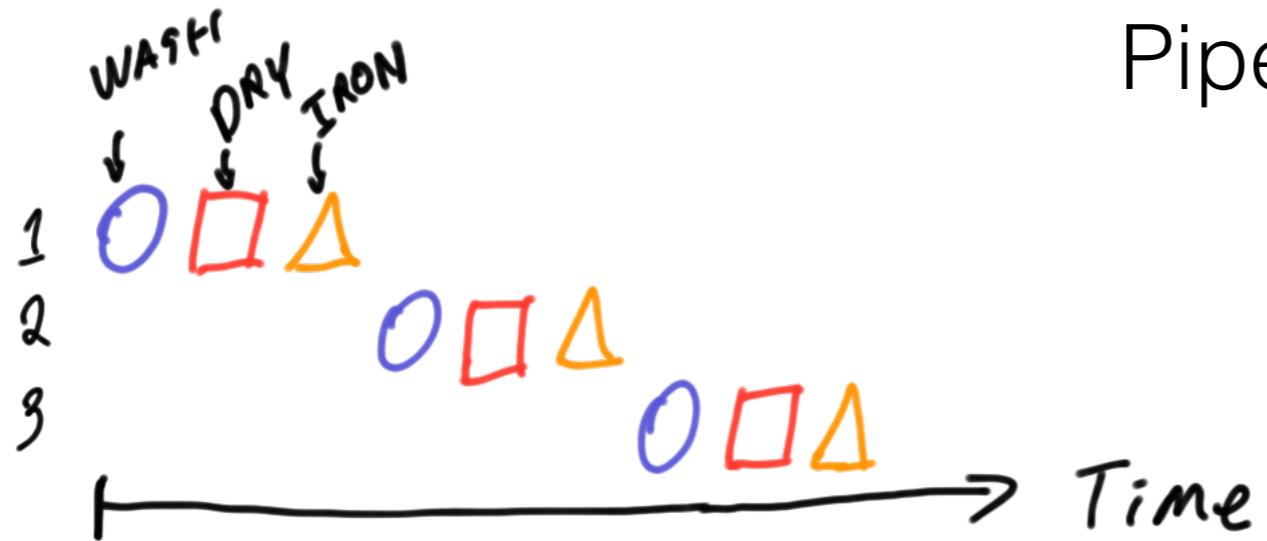


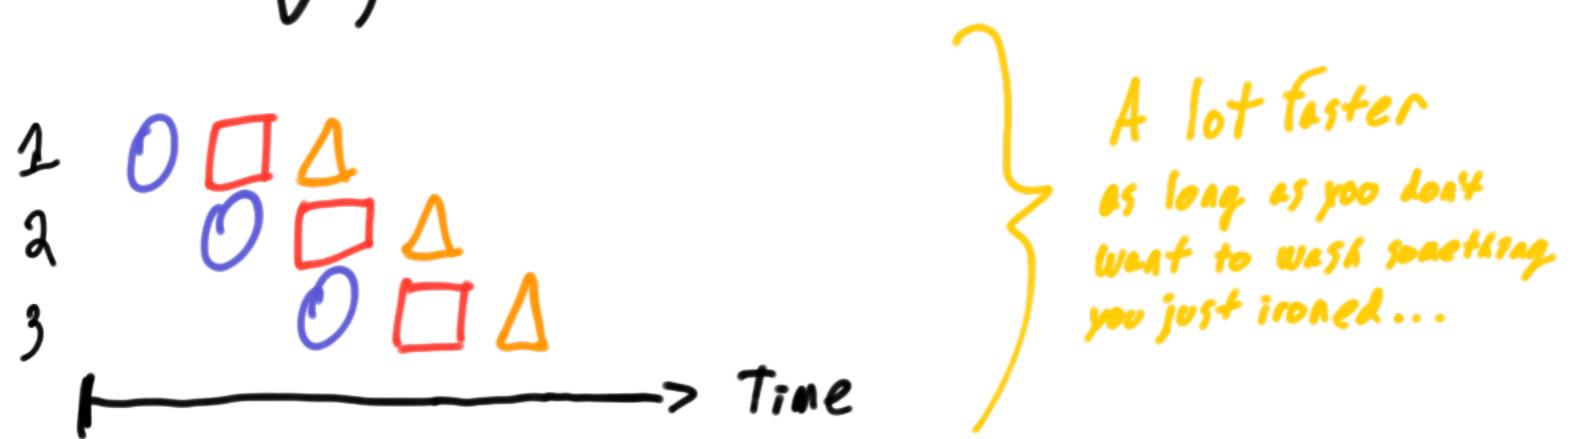
FIGURE 4.2 The basic implementation of the MIPS subset, including the necessary multiplexors and control

The top multiplexor ("Mux") controls what value replaces the PC ($PC + 4$ or the branch destination address); the multiplexor is controlled by the gate that "ANDs" together the Zero output of the ALU and a control signal that indicates that the instruction is a branch. The middle multiplexor, whose output returns to the register file, is used to steer the output of the ALU (in the case of an arithmetic-logical instruction) or the output of the data memory (in the case of a load) for writing into the register file. Finally, the bottommost multiplexor is used to determine whether the second ALU input is from the registers (for an arithmetic-logical instruction or a branch) or from the offset field of the instruction (for a load or store). The added control lines are straightforward and determine the operation performed at the ALU, whether the data memory should read or write, and whether the registers should perform a write operation. The control lines are shown in color to make them easier to see.

Pipelining



VS



Pipelining

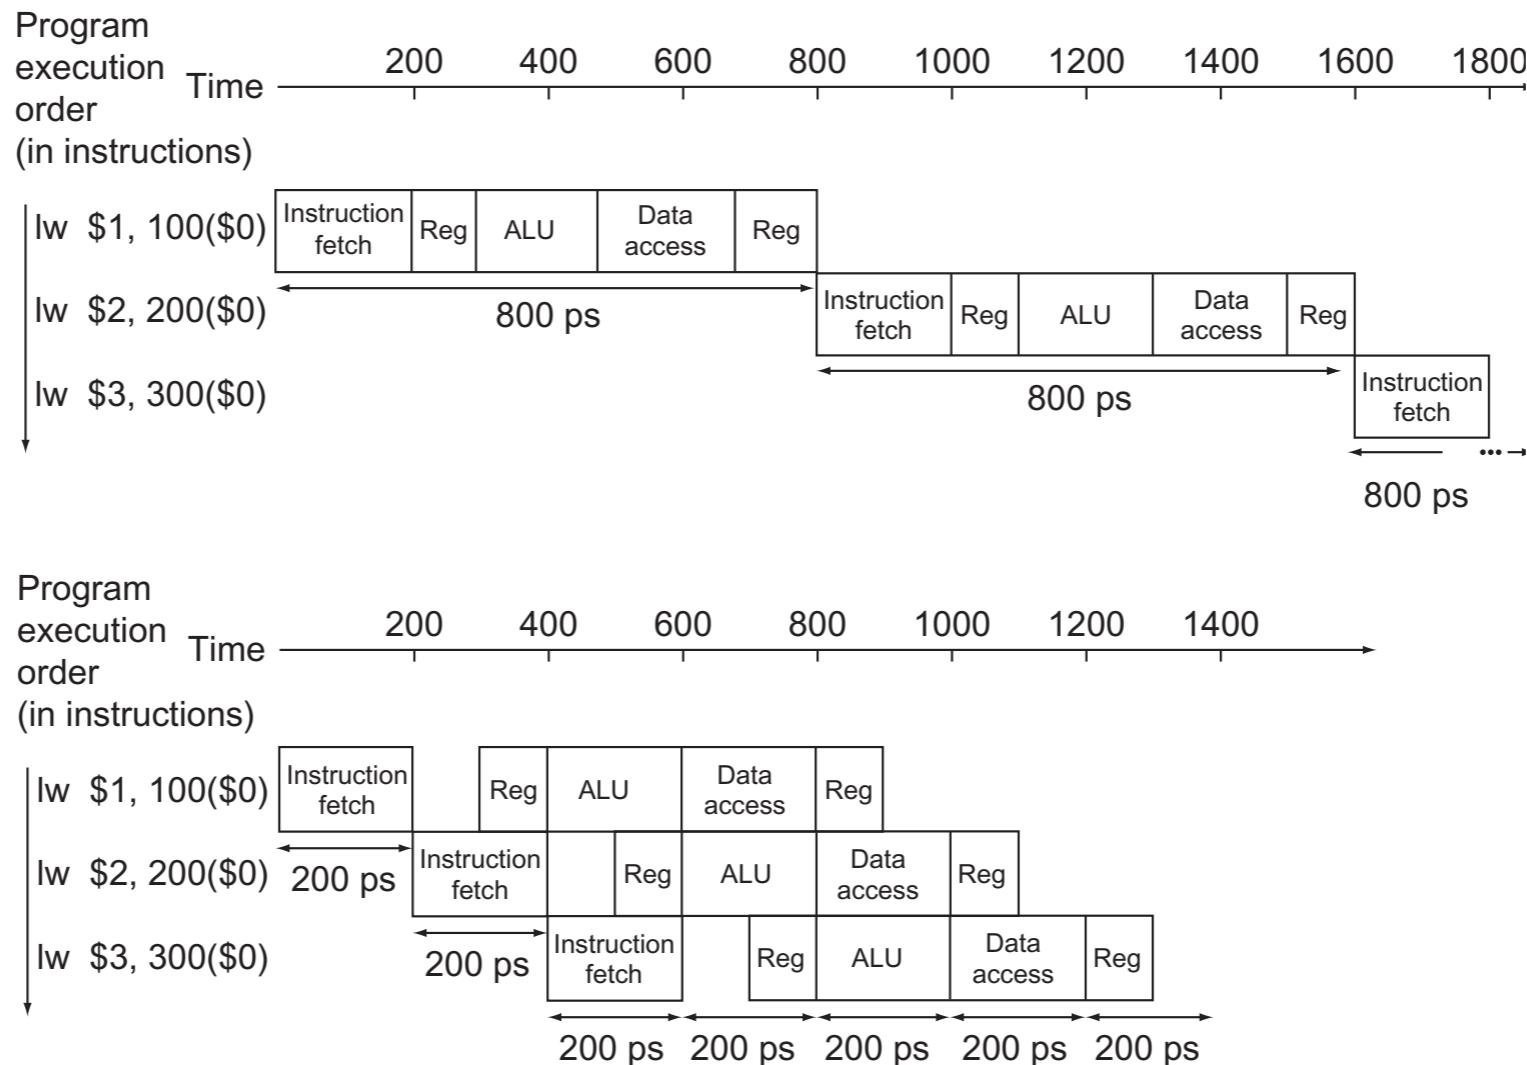


FIGURE 4.27 Single-cycle, nonpipelined execution in top versus pipelined execution in bottom. Both use the same hardware components, whose time is listed in Figure 4.26. In this case, we see a fourfold speed-up on average time between instructions, from 800 ps down to 200 ps. Compare this figure to Figure 4.25. For the laundry, we assumed all stages were equal. If the dryer were slowest, then the drye stage would set the stage time. The pipeline stage times of a computer are also limited by the slowest resource either the ALU operation or the memory access. We assume the write to the register file occurs in the first half of the clock cycle and the read from the register file occurs in the second half. We use this assumption throughout this chapter.

Latency

"Latency numbers every programmer should know"

↙ One CPU clock cycle
↳ (so must be for a 86Hz CPU)

L1 cache reference	0.5 ns			
Branch mispredict	5 ns			
L2 cache reference	7 ns			14x L1 cache
Main memory reference	100 ns			200x L1 cache
Read 1 MB sequentially from memory	250,000 ns	250 us		
Read 1 MB sequentially from SSD*	1,000,000 ns	1,000 us	1 ms	
Disk seek	10,000,000 ns	10,000 us	10 ms	
Read 1 MB sequentially from disk	20,000,000 ns	20,000 us	20 ms	

So tl;dr: RAM is 200 clock cycles away
HD is 20,000,000 clock cycles away