


Technical Computing for the Earth
Sciences, Lecture 10:

Introduction to Parallel Programming

EARS 80.03

CPU basics



CPU basics

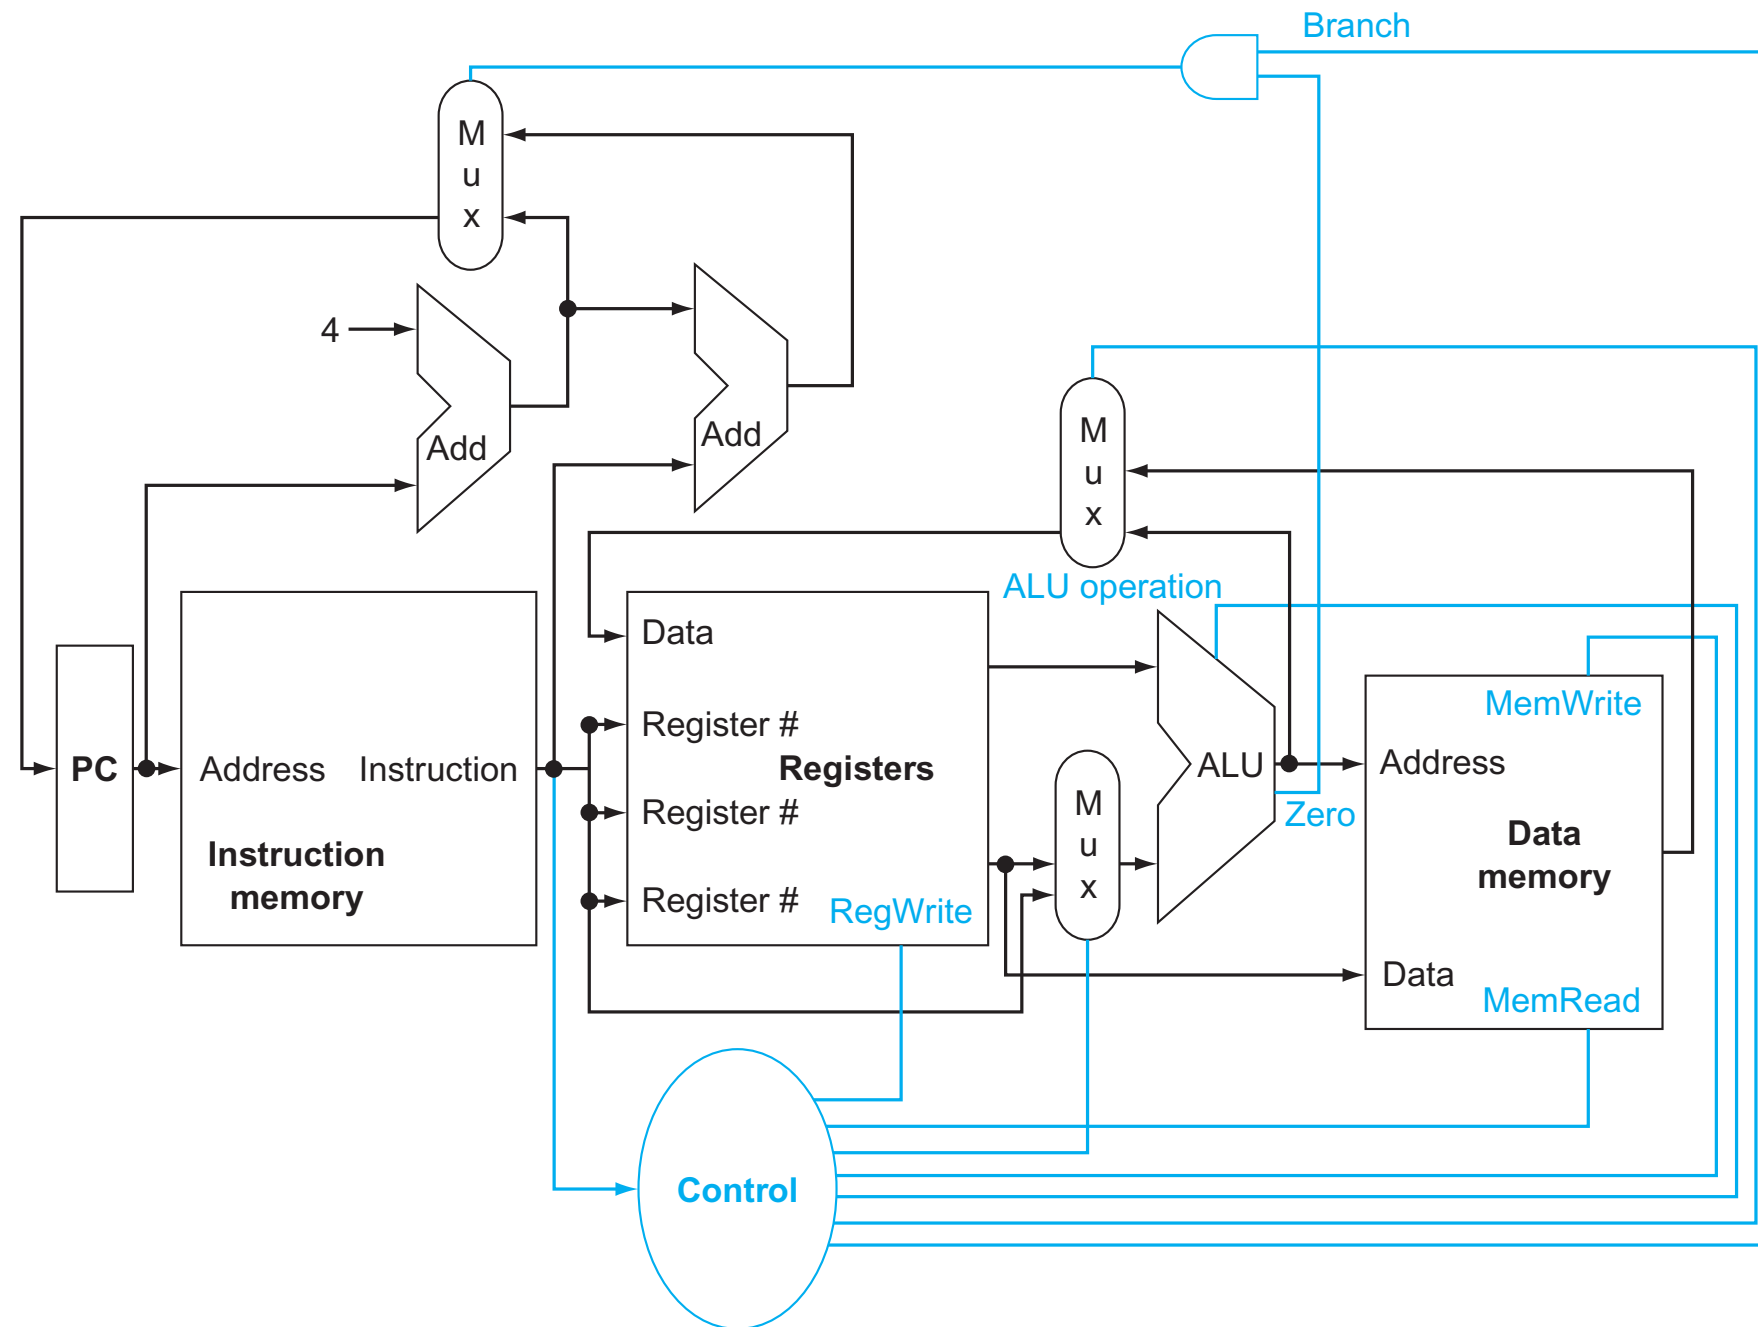



FIGURE 4.2 The basic implementation of the MIPS subset, including the necessary multiplexors and control

The top multiplexor (“Mux”) controls what value replaces the PC ($PC + 4$ or the branch destination address); the multiplexor is controlled by the gate that “ANDs” together the Zero output of the ALU and a control signal that indicates that the instruction is a branch. The middle multiplexor, whose output returns to the register file, is used to steer the output of the ALU (in the case of an arithmetic-logical instruction) or the output of the data memory (in the case of a load) for writing into the register file. Finally, the bottommost multiplexor is used to determine whether the second ALU input is from the registers (for an arithmetic-logical instruction or a branch) or from the offset field of the instruction (for a load or store). The added control lines are straightforward and determine the operation performed at the ALU, whether the data memory should read or write, and whether the registers should perform a write operation. The control lines are shown in color to make them easier to see.

Pipelining



VS



Pipelining

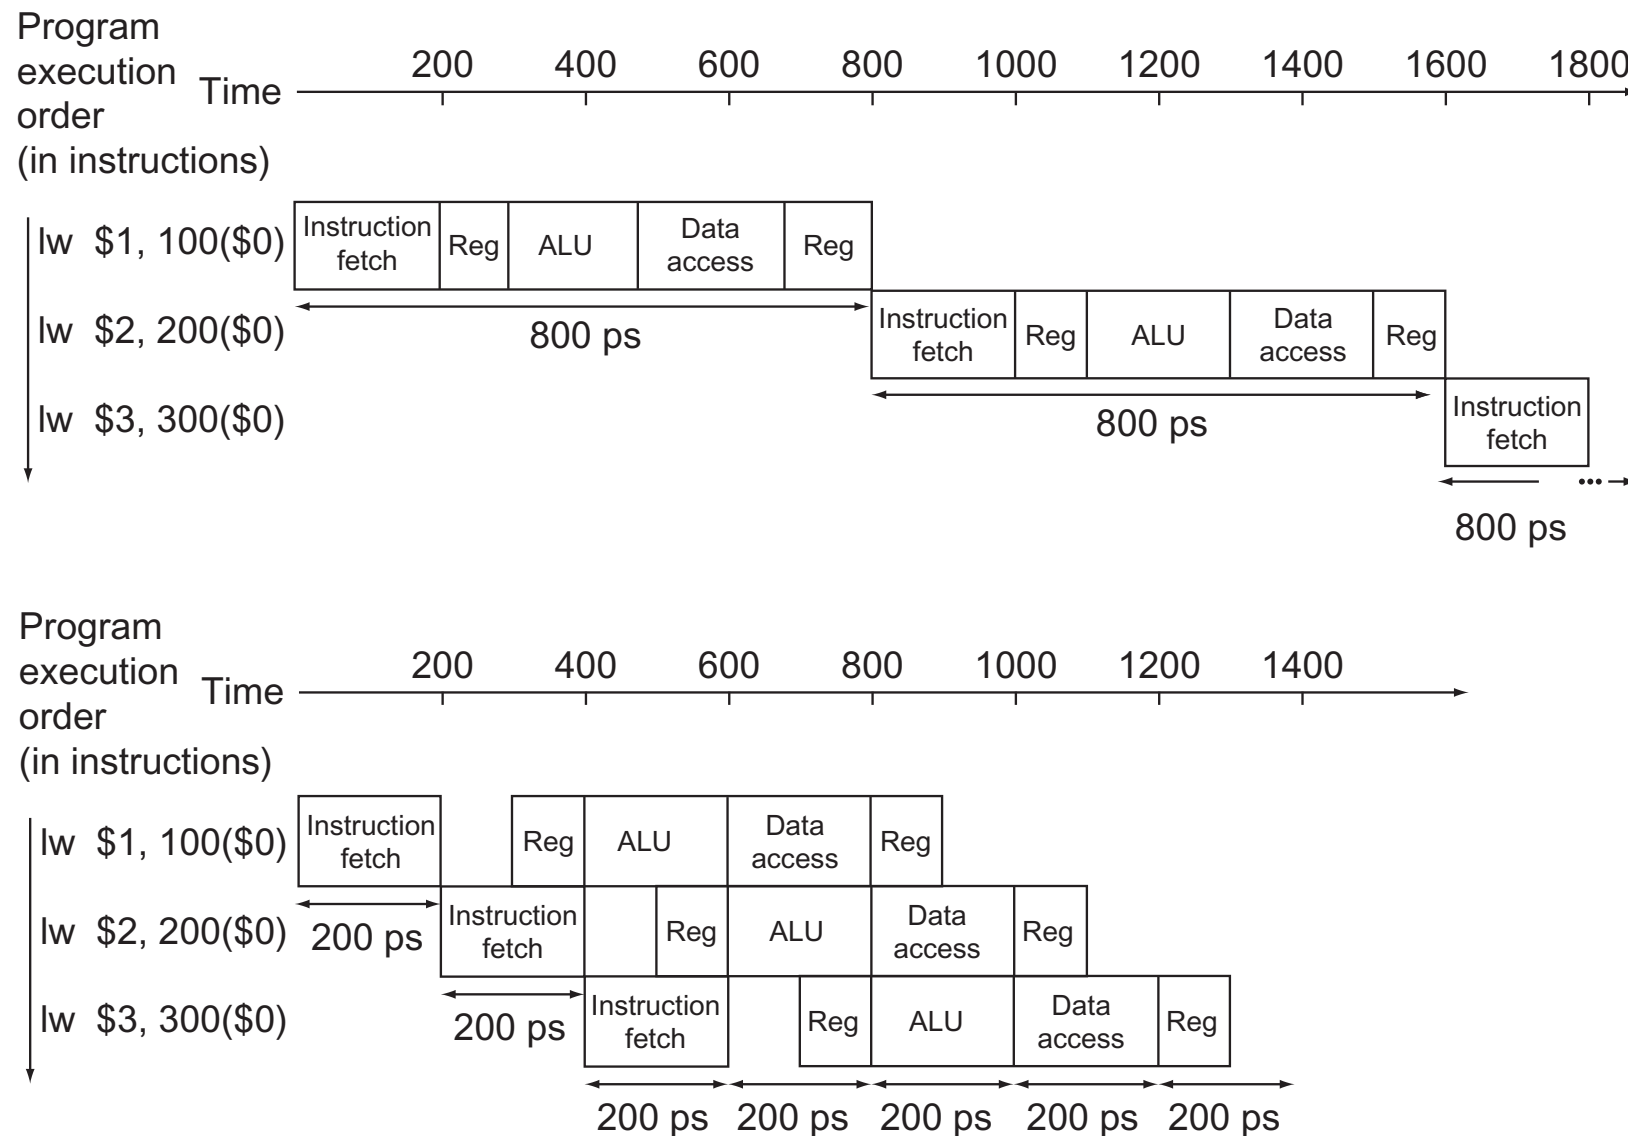
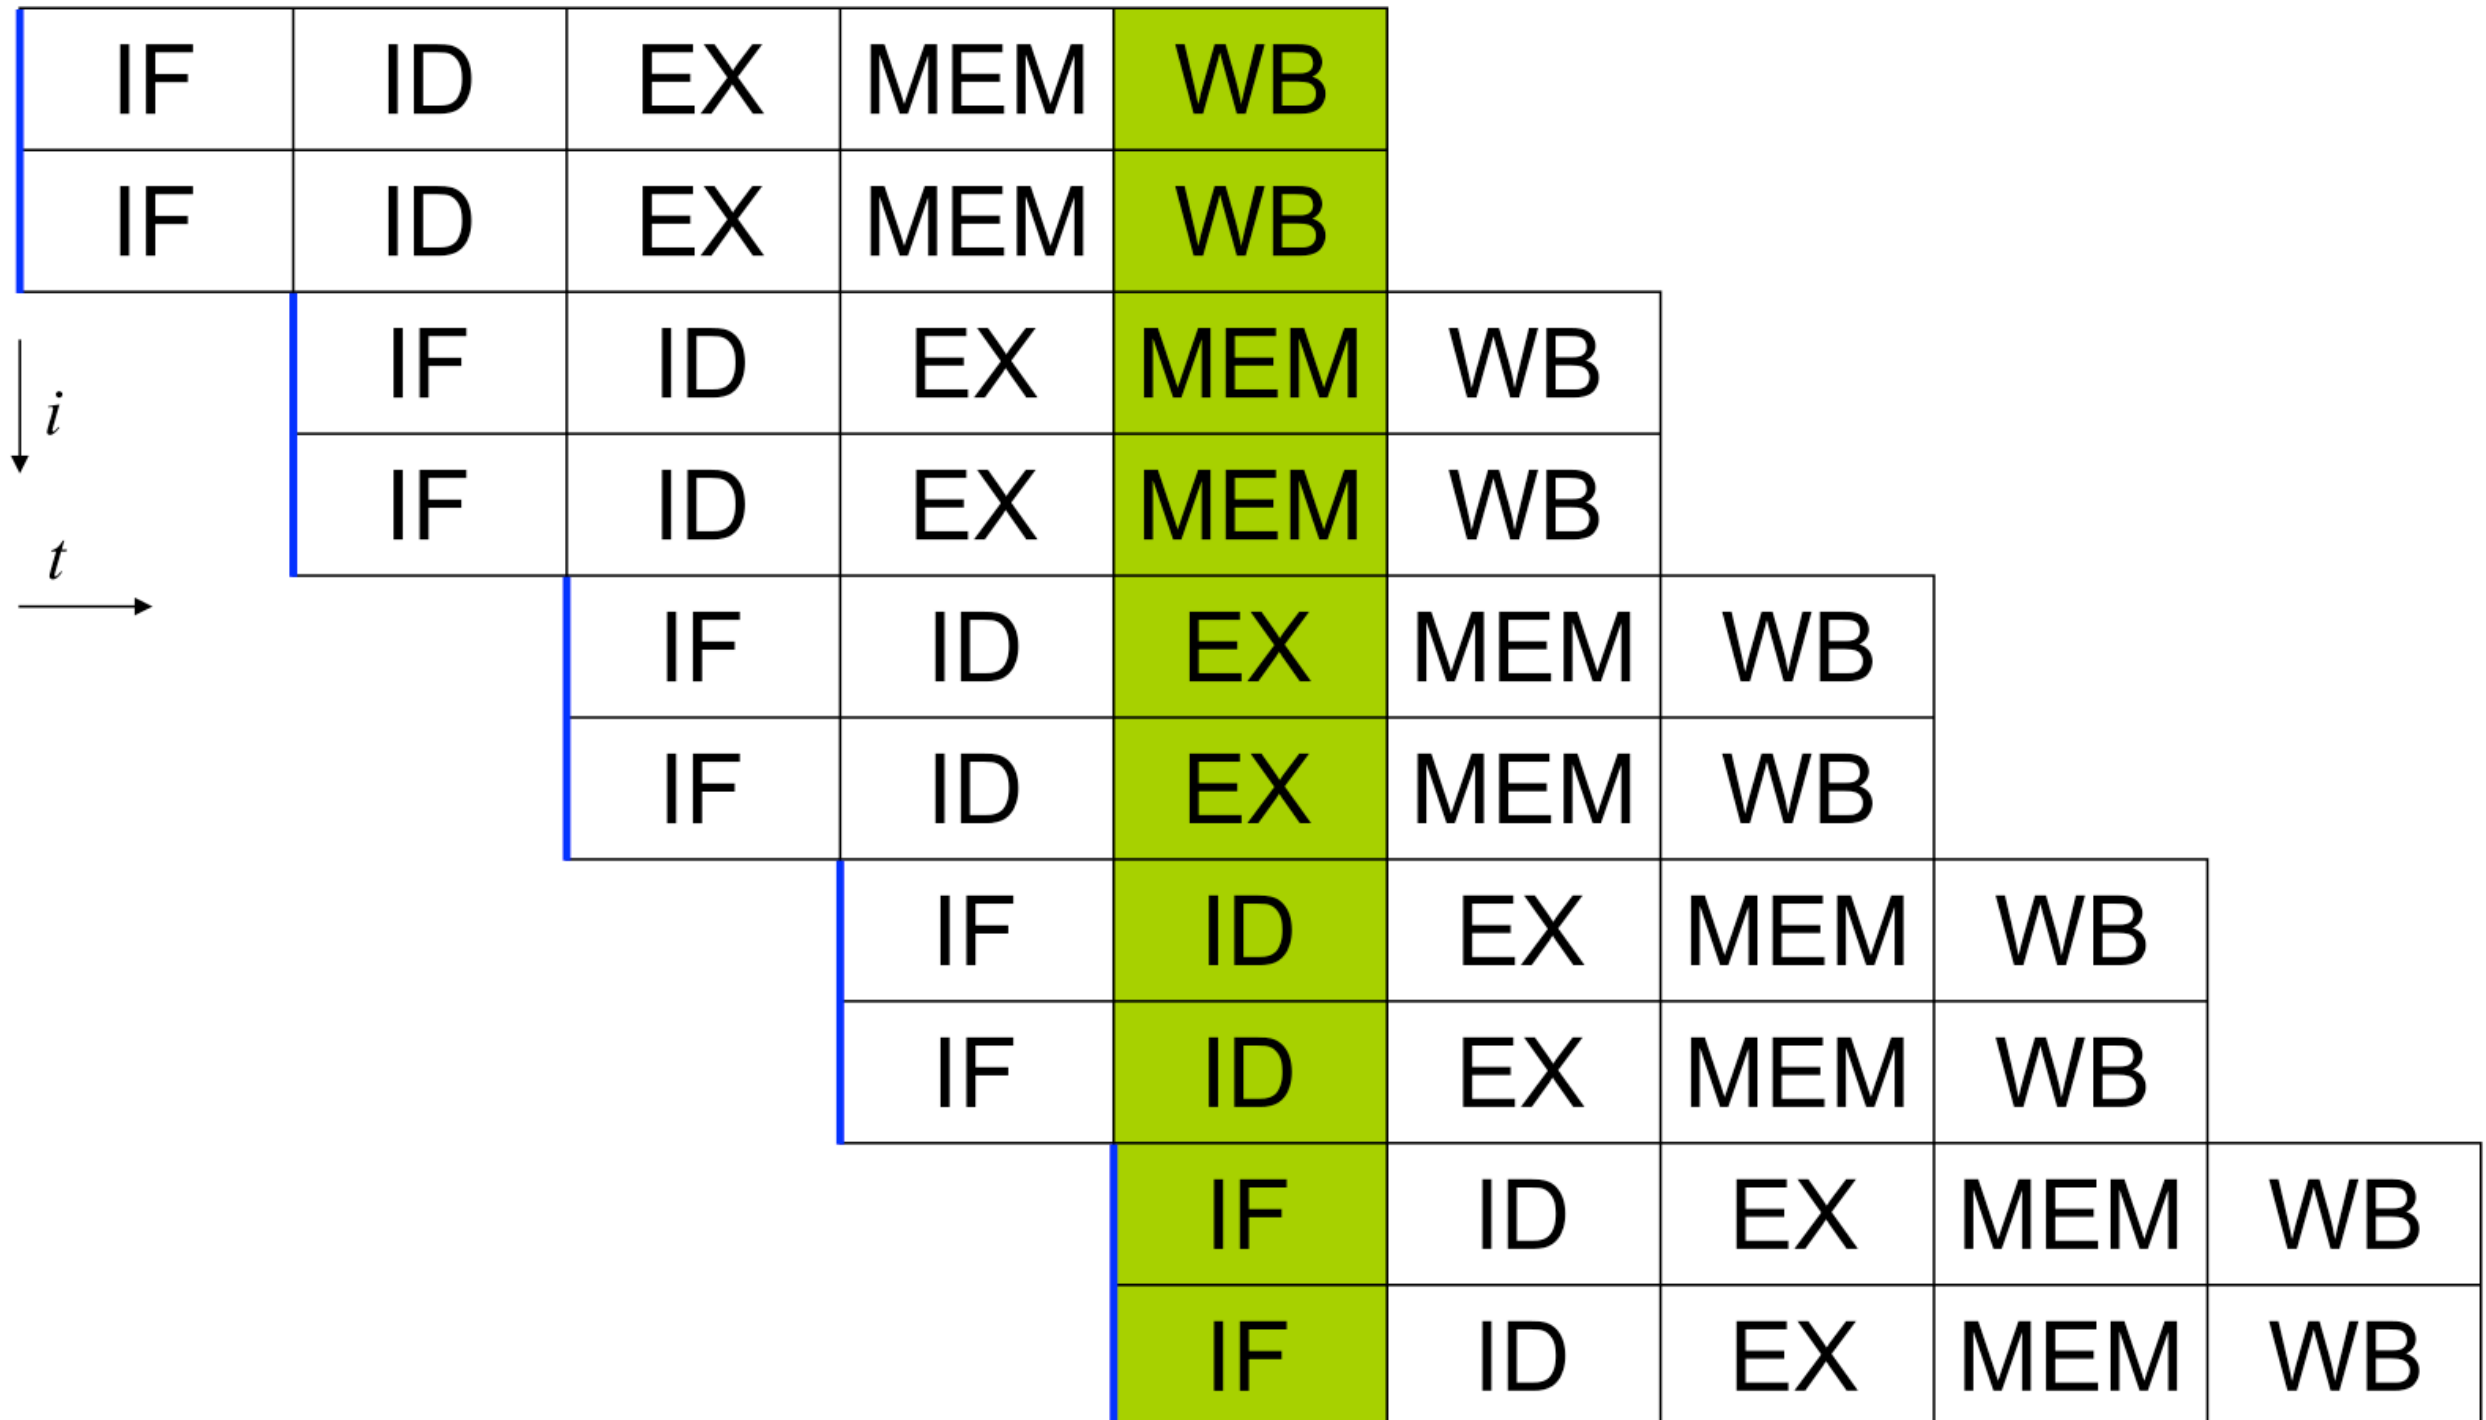
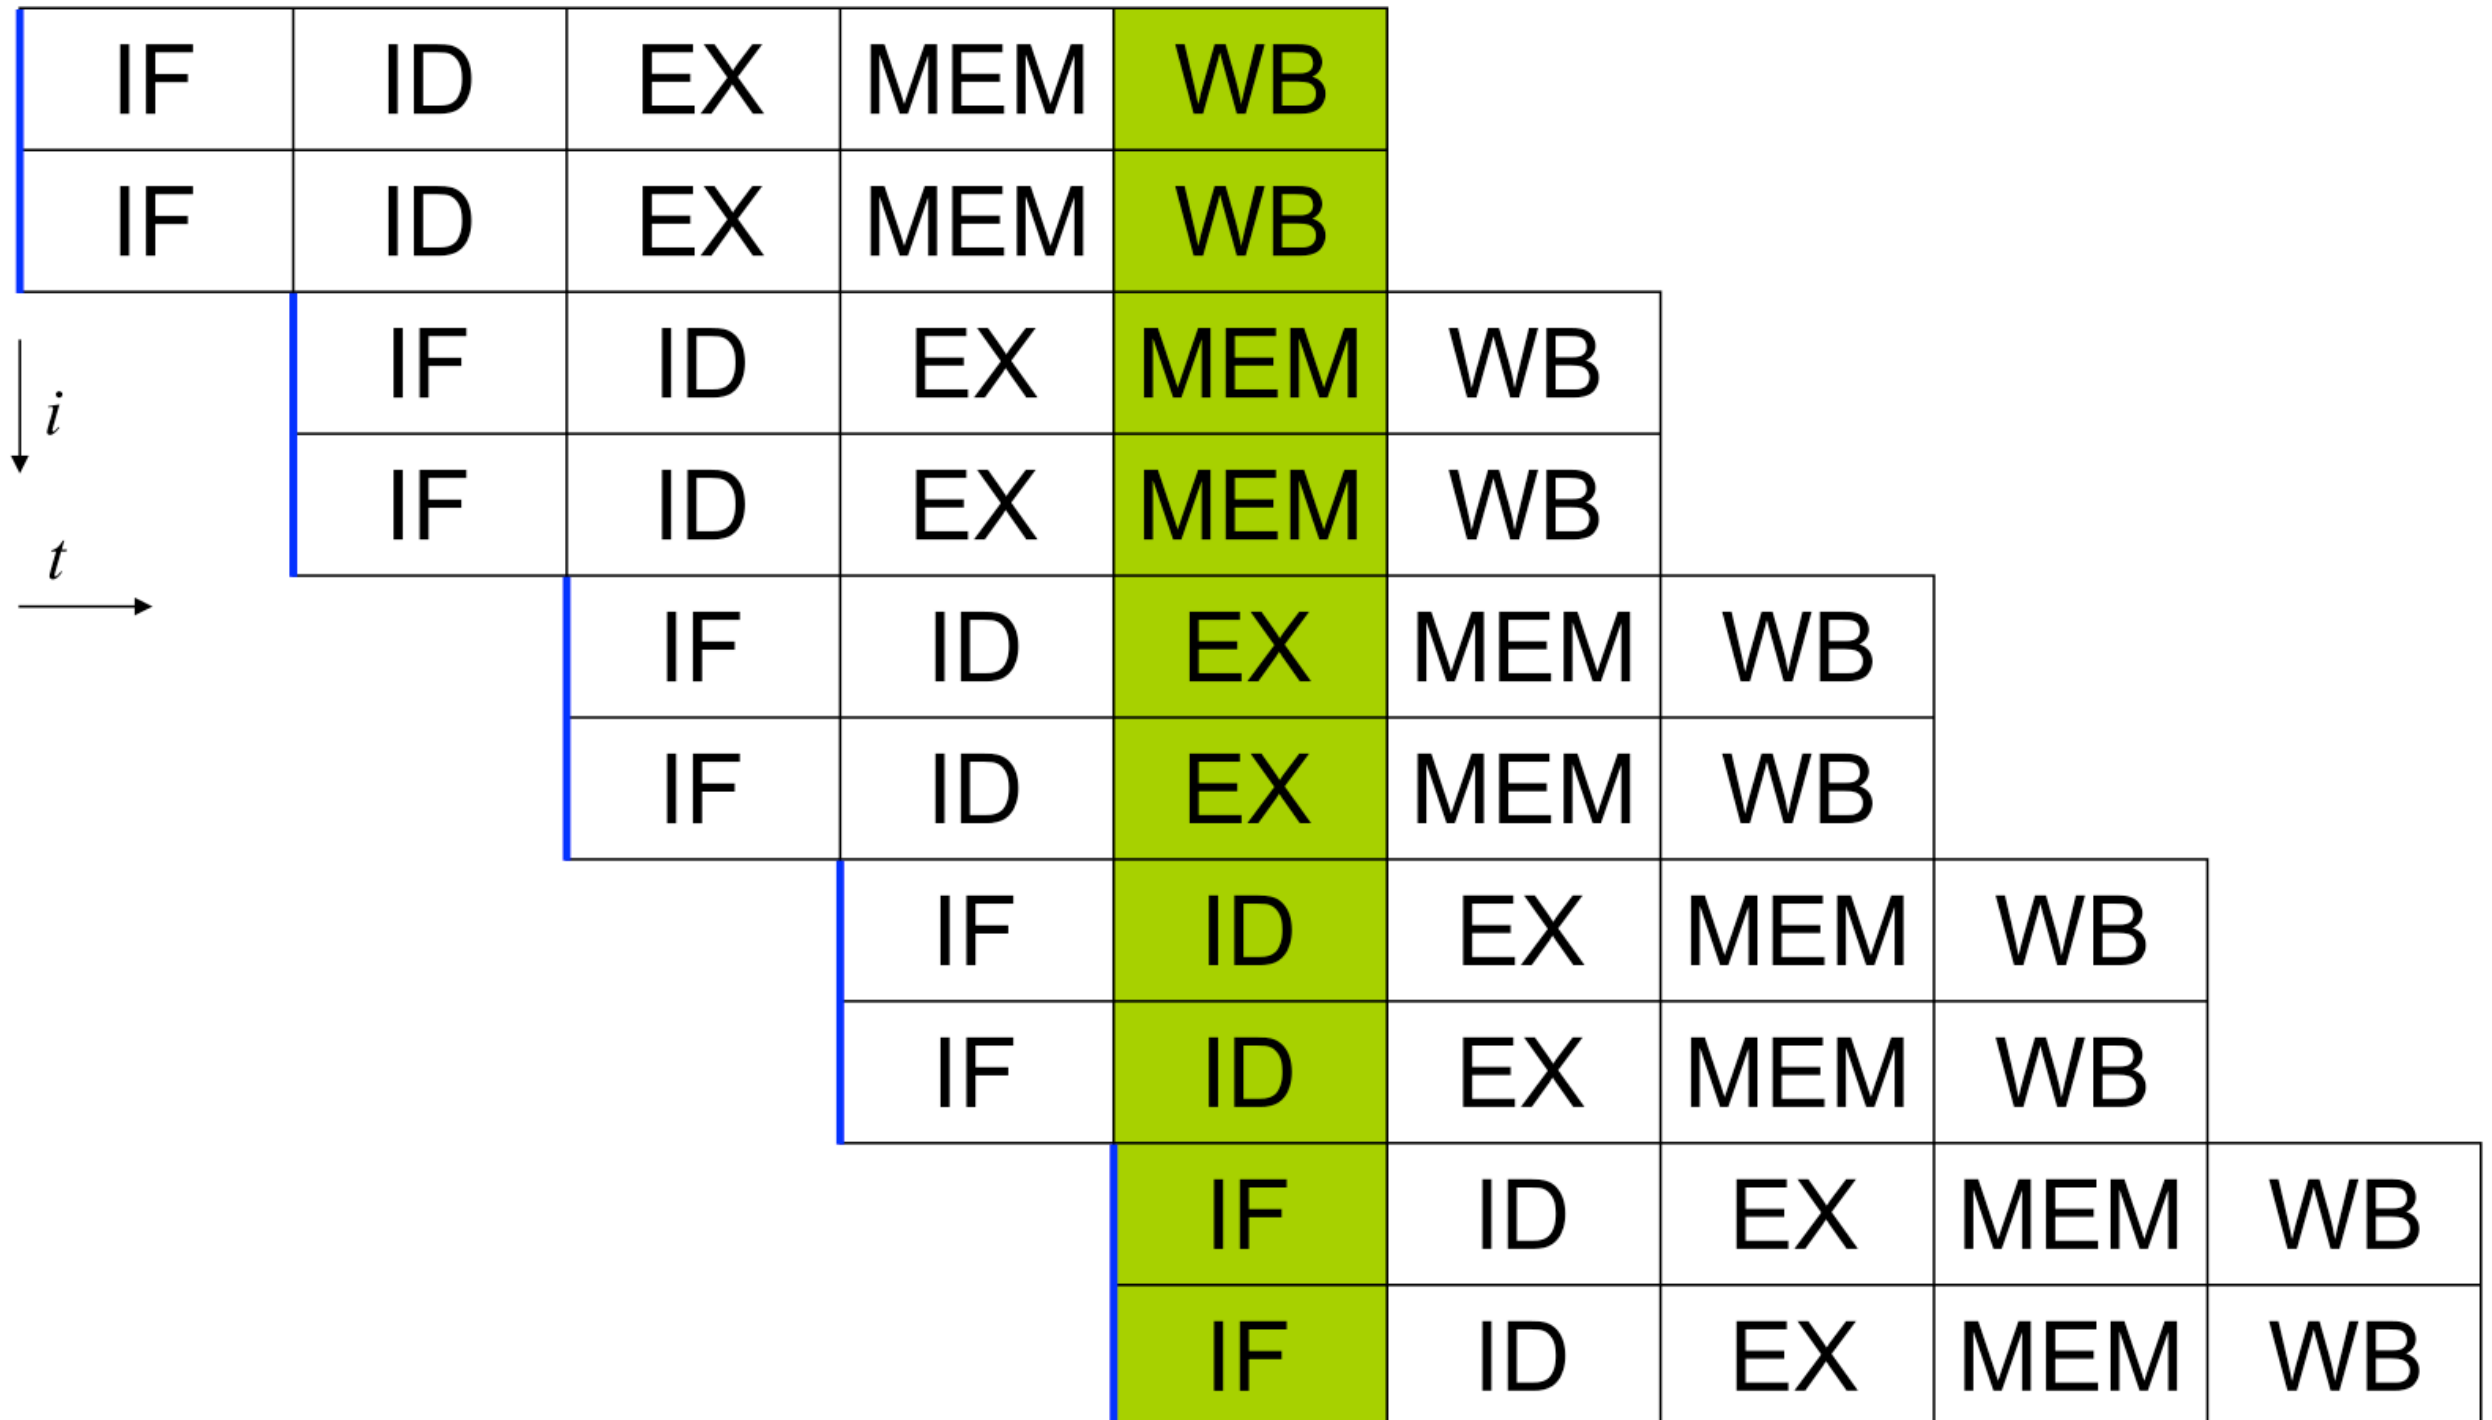


FIGURE 4.27 Single-cycle, nonpipelined execution in top versus pipelined execution in bottom. Both use the same hardware components, whose time is listed in [Figure 4.26](#). In this case, we see a fourfold speed-up on average time between instructions, from 800 ps down to 200 ps. Compare this figure to [Figure 4.25](#). For the laundry, we assumed all stages were equal. If the dryer were slowest, then the dryer stage would set the stage time. The pipeline stage times of a computer are also limited by the slowest resource either the ALU operation or the memory access. We assume the write to the register file occurs in the first half of the clock cycle and the read from the register file occurs in the second half. We use this assumption throughout this chapter.

Actually...

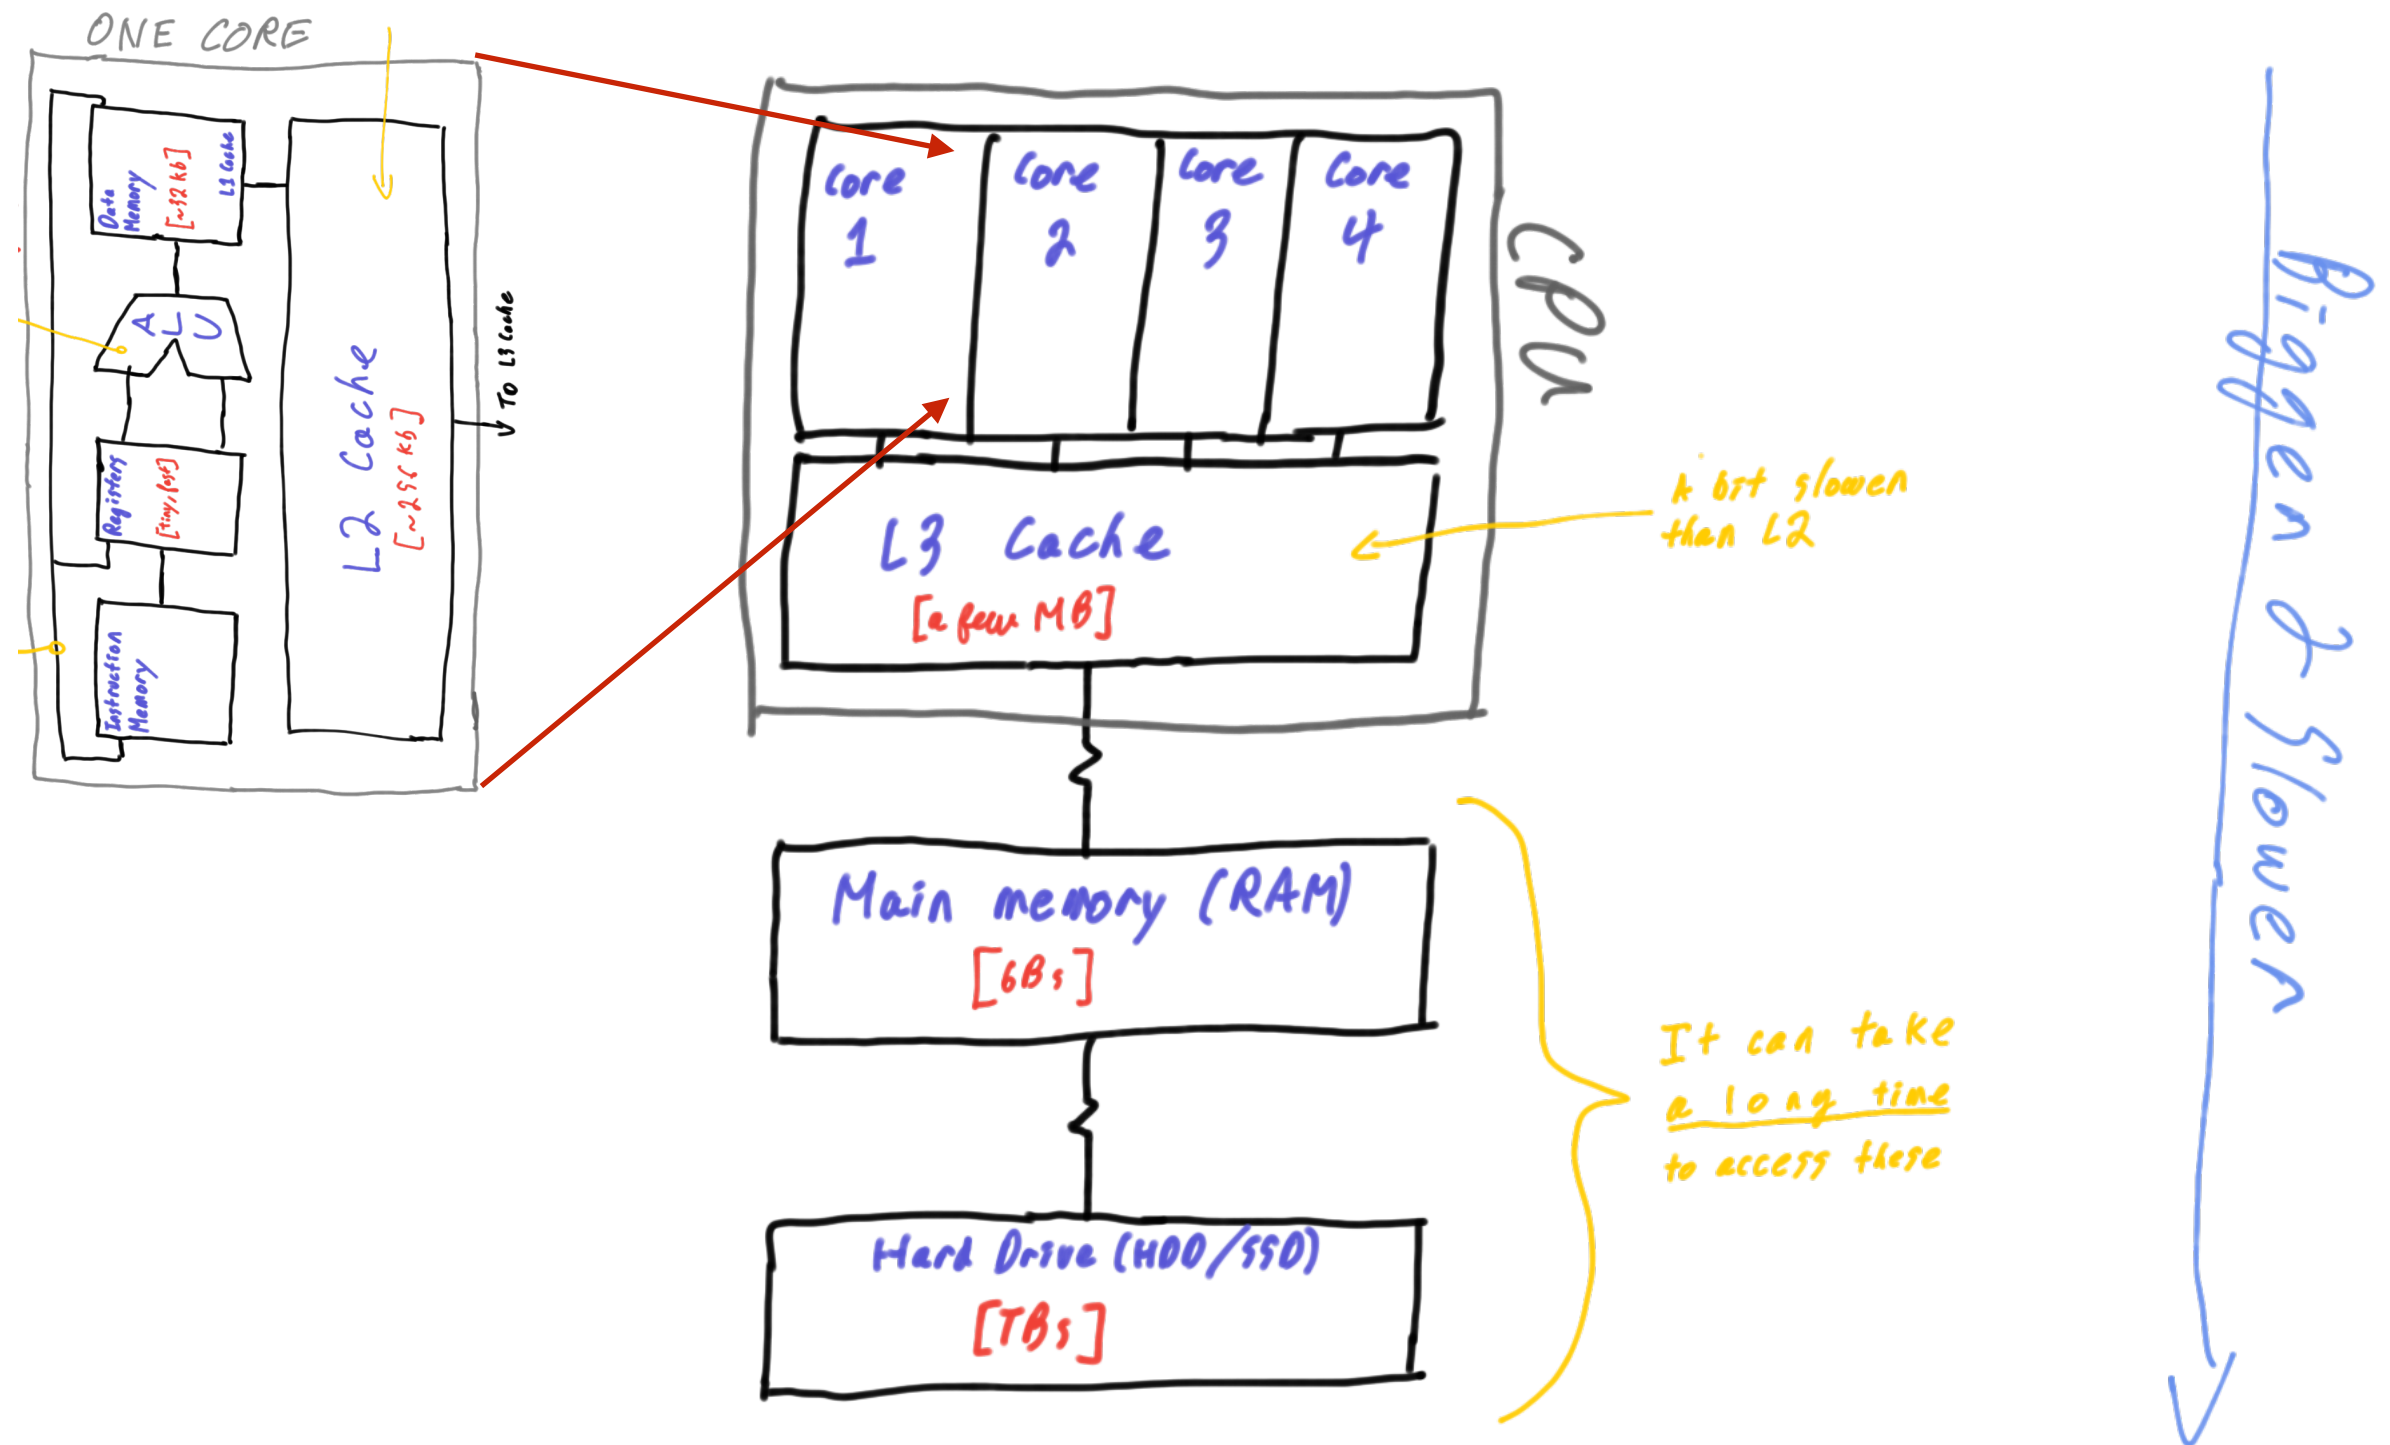


Actually...




Solution: LoopVectorization.jl and @avx


Meanwhile, there are now multiple cores in a CPU



Meanwhile, there are now multiple cores in a CPU



Meanwhile, there are now multiple cores in a CPU



Solution: Shared-memory parallel programming
e.g. Base.Threads.@threads, etc.

Latency

"Latency numbers every programmer should know"

L1 cache reference		0.5	ns		
Branch mispredict		5	ns		
L2 cache reference		7	ns		
Main memory reference		100	ns		
Read 1 MB sequentially from memory	250,000	ns	250	us	
Read 1 MB sequentially from SSD*	1,000,000	ns	1,000	us	1 ms
Disk seek	10,000,000	ns	10,000	us	10 ms
Read 1 MB sequentially from disk	20,000,000	ns	20,000	us	20 ms

14x L1 cache
200x L1 cache